



Big Data Management and Analytics

Session: Distributed Storage - Hadoop III

Lecturers: Petar Jovanovic and Josep Berbegal

1 Tasks To Do Before The Session

It is important that you: (1) carefully read the **instruction sheet** for this lab session, (2) introduce yourself to the **lab's main objectives**, (3) understand the **theoretical background**, and (4) get familiar with the **tools being used**.

2 Part A: Examples & Questions (15min)

In the first 15 minutes, we will first clarify the main objectives of this lab. We will then learn how HDFS stores your data in the disk and overview different possible formats to do so, including their benefits and overhead. Finally, we will also see how HDFS can compress your data to optimize the storage space.

3 Part B: In-class Practice (2h 45min)

In this exercise you will use the provided Java program in order to generate data with different formats. The run configuration arguments are the following (in the same order): *format number_of_tuples path_for_file*. Precisely, we will explore the following formats (in bold the argument that will be used in the run configuration):

1. Plain text (-plainText).
2. Sequence files (-sequenceFile).
3. Avro (-avro).
4. Parquet (-parquet).



3.1 Exercise 1 (1h): Writing and Reading in different formats

1. Generate files with 10000 rows as a plain text, and then as well with Sequence files, Avro, and Parquet formats, and insert them into HDFS.
Run:

```
java -jar Hadoop-HDFS-0.1.jar write -plainText 10000 wines.10000.txt
java -jar Hadoop-HDFS-0.1.jar write -sequenceFile 10000 wines.10000.seq
java -jar Hadoop-HDFS-0.1.jar write -avro 10000 wines.10000.avr
java -jar Hadoop-HDFS-0.1.jar write -parquet 10000 wines.10000.prq
```

Then, load these files into your HDFS using the put command.

2. Read each of your files, with the following standard commands (-cat):

```
hadoop-2.7.4/bin/hdfs dfs -cat wines.10000.txt
hadoop-2.7.4/bin/hdfs dfs -cat wines.10000.seq
hadoop-2.7.4/bin/hdfs dfs -cat wines.10000.avr
hadoop-2.7.4/bin/hdfs dfs -cat wines.10000.prq
```

- What do your files look like now? Are they readable?

Answer:

- What do you think is happening?

Answer:

3. Now, read the file stored in the SequenceFile format again by means of the following command:

```
java -jar Hadoop-HDFS-0.1.jar read -sequenceFile wines.10000.seq
```

- Is your file readable now?

Answer:

- What do you think the first column represents?

Answer:

- Is it something given automatically or do we need to specify it?

Answer:



- **How is it built?**

Hint: Follow the code at the classes `MyHDFSSequenceFileWriter` and `MyHDFSSequenceFileReader` for answers.

Answer:

3.2 Exercise 2 (45min): Format overhead

3.2.1 Exercise 2.1

Now, generate 1 tuple for each format. To this end, use the appropriate run configuration for each. For instance, you would use the following parameters for the plain text case (the rest of the command is like in the previous cases): `write -plainText 1 winesplainText.txt`.

Next, answer the following questions:

- (a) **What is the exact size for each file?**

Answer:

- (b) **For Sequence files, Avro, and Parquet what is the overhead ratio with respect to plain text (i.e., the result of the fraction $\text{sizeOfFormatX}/\text{sizeOfPlainText}$)?**

Answer:

- (c) **Why? Elaborate on why you observe this file size with respect to the properties of the used format.**

Answer:

3.2.2 Exercise 2.2

Second, generate 1000 tuples for each format and answer the following questions:

- (a) **What is the exact size for each file?**

Answer:



- (b) For Sequence files, Avro and Parquet what is the fraction of overhead with respect to plain text?

Answer:

- (c) If the overhead ratio is different, explain why.

Answer:

3.3 Exercise 3 (1h): Data compression

Fill table 1. Here you need to change (actually uncomment) the code at the class *MyHDFSSequenceFileWriter*¹ (remember to recompile and redeploy it after the change) in order to write a sequence file of 2 million rows but by using different compression levels (line 43 - NONE, line 44 - RECORD, line 45 - BLOCK). Initially, it is set to NONE! Importantly, remember to recompile and redeploy the project, following the instructions in the first session, and produce new JAR file (for convenience you can rename it by adding the corresponding suffix: '-NONE', '-RECORD', '-BLOCK').

Finally explain the meaning of each compression level and discuss about the results you obtain.

Level	Insertion time	Reading time	#Blocks	File size
NONE				
RECORD				
BLOCK				

Table 1: Compression on sequence files

Reasoning on Table 1

¹.../src/main/java/bdma/labos/hadoop/writer/MyHDFSSequenceFileWriter.java



UNIVERSITAT POLITÈCNICA DE CATALUNYA
BARCELONATECH

School of Professional & Executive Development

Additional comments: